

Scoreboard Summary

- Speedup 1.7 from compiled code; 2.5 for hand-coded
- Limitations of 6600 scoreboard:
 - No forwarding hardware
 - Limited to instructions in basic block (small *window*)
 - why?
 - Small number of functional units (structural hazards)
 - insts to same fu cannot be reordered
 - Wait for WAR hazards (after EX, before WB)
 - Prevent WAW hazards (in ID)

CSE 240A

Dean Tullsen

Another Dynamic Algorithm: Tomasulo Algorithm

- For IBM 360/91 about 3 years after CDC 6600
- Goal: High Performance without special compilers
- Differences between IBM 360 & CDC 6600 ISA
 - IBM has only 2 register specifiers/instr vs. 3 in CDC 6600
 - IBM has 4 FP registers vs. 8 in CDC 6600
 - Implications?

CSE 240A

Dean Tullsen

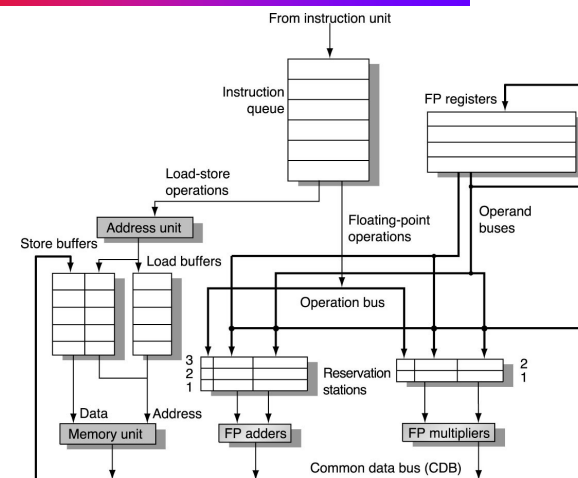
Differences between Tomasulo Algorithm & Scoreboard

- Control & buffers distributed with Function Units vs. centralized in scoreboard; called “reservation stations”
 - => instrs schedule themselves
- Registers in instructions replaced by pointers to reservation station buffer scoreboard => registers primary operand storage
Tomasulo => reservation stations as operand storage
- HW renaming of registers to avoid WAR, WAW hazards
Scoreboard => both source registers read together (thus one could not be overwritten while we wait for the other).
Tomasulo => each register read as soon as available.
- Common Data Bus broadcasts results to all FUs
RS's (FU's), registers, etc. responsible for collecting own data off CDB
- Load and Store Queues treated as FUs as well

CSE 240A

Dean Tullsen

Tomasulo Organization



CSE 240A

Dean Tullsen

Reservation Station Components

Op—Operation to perform in the unit (e.g., + or −)

Qj, Qk—Reservation stations producing source registers

Vj, Vk—Value of Source operands

Rj, Rk—Flags indicating when Vj, Vk are ready

Busy—Indicates reservation station is busy

Register result status—Indicates which functional unit will write each register, if one exists. Blank when no pending instructions that will write that register.

Three Stages of Tomasulo Algorithm

1. Issue—get instruction from FP Op Queue

If reservation station free, the scoreboard issues instr & sends operands (renames registers).

2. Execution—operate on operands (EX)

When both operands ready then execute;
if not ready, watch CDB for result

3. Write result—finish execution (WB)

Write on Common Data Bus to all waiting units;
mark reservation station available.

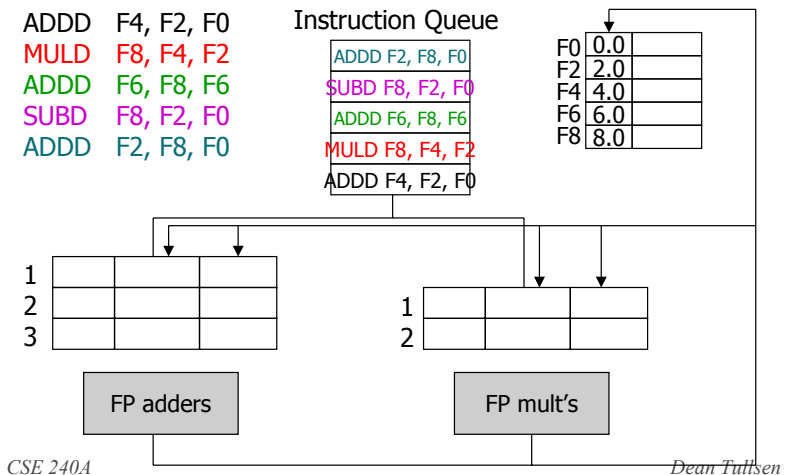
Tomasulo Example

```

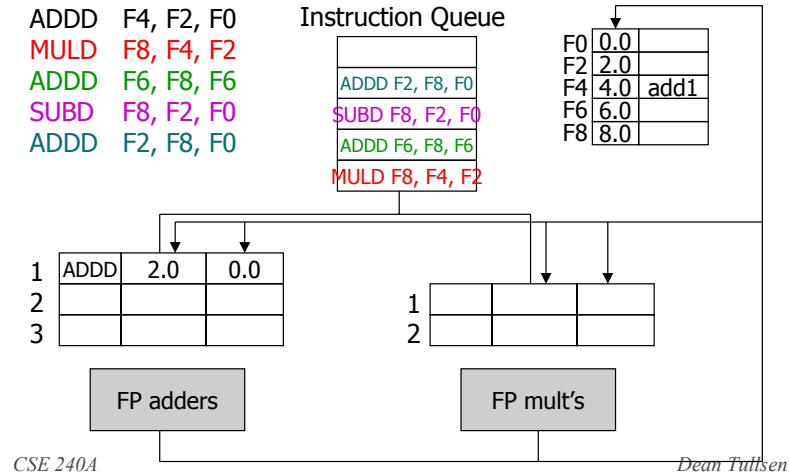
ADDD  F4, F2, F0
MULD  F8, F4, F2
ADDD  F6, F8, F6
SUBD  F8, F2, F0
ADDD  F2, F8, F0
    
```

Multiply takes 10 clocks, add/sub take 4

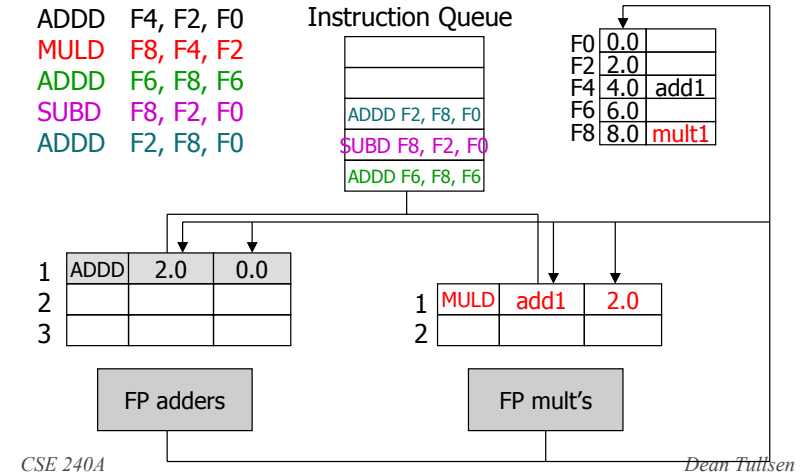
Tomasulo – cycle 0



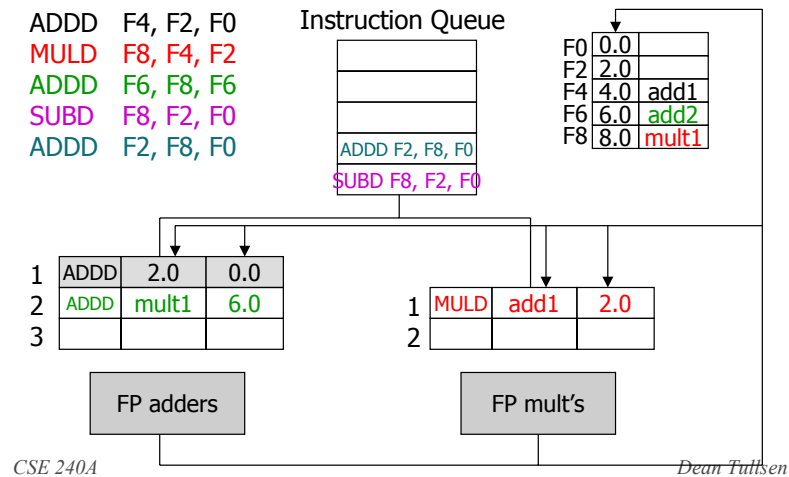
Tomasulo – cycle 1



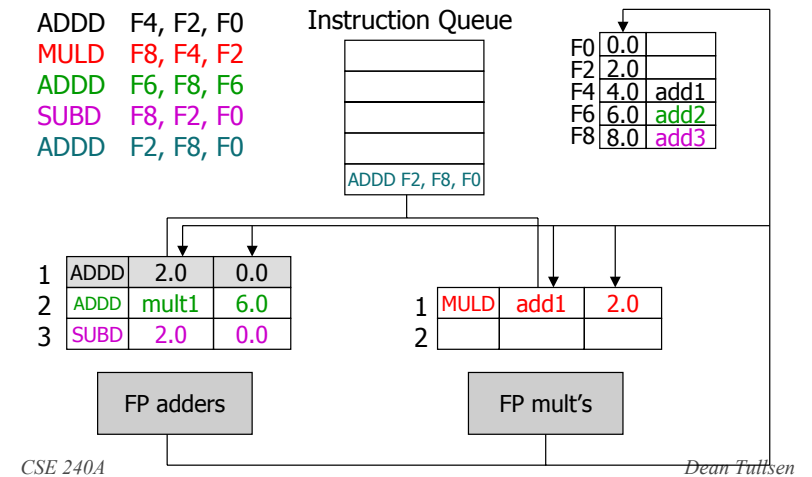
Tomasulo – cycle 2



Tomasulo – cycle 3



Tomasulo – cycle 4





Dean Tuttsen



Dean Tullsen

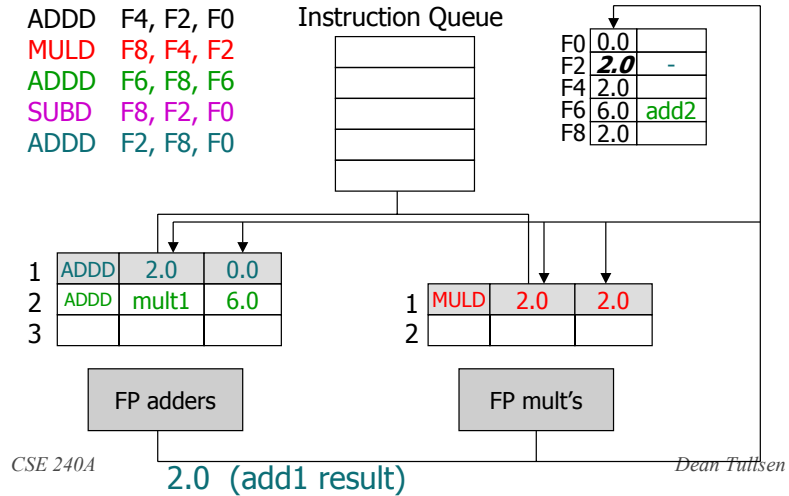


~~Dean Tuttsen~~

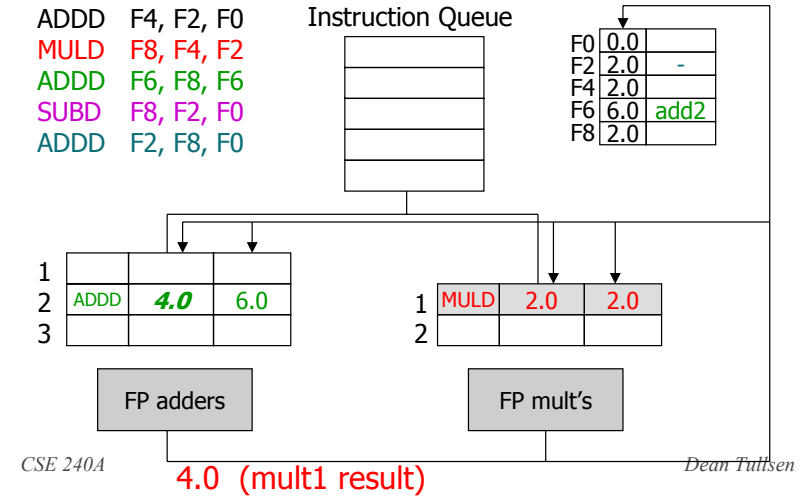


~~Dean Tuttsen~~

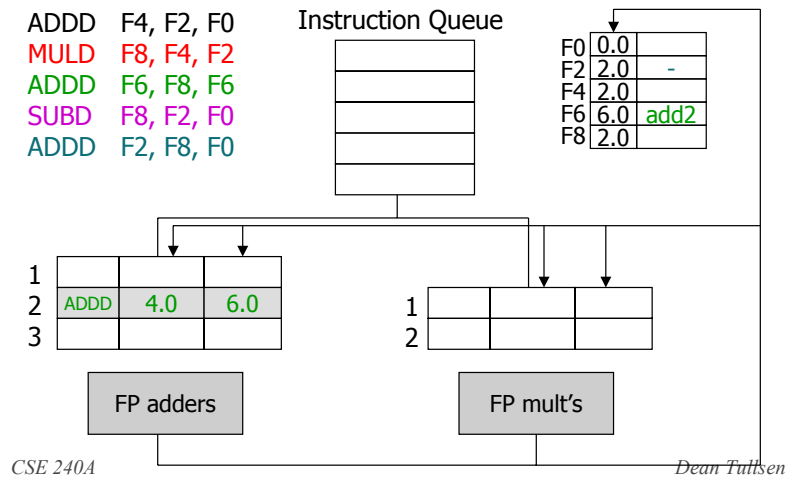
Tomasulo - cycle 12



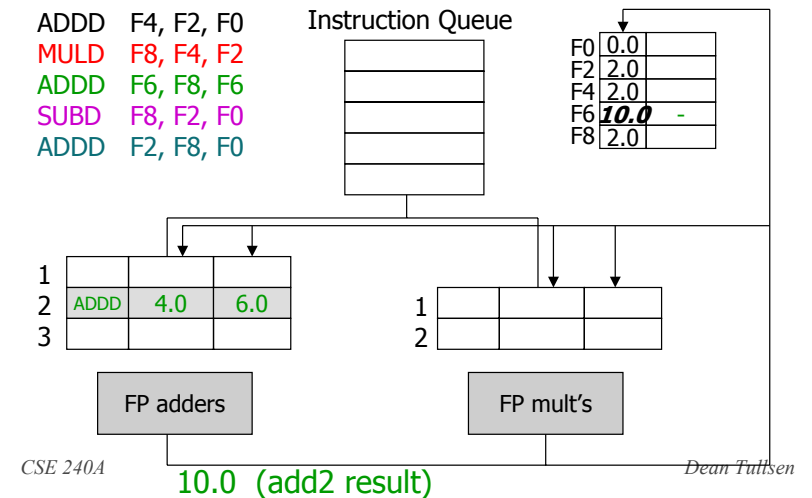
Tomasulo - cycle 15



Tomasulo - cycle 16



Tomasulo - cycle 19



Tomasulo Summary

- Prevents Register as bottleneck
- Avoids WAR, WAW hazards of Scoreboard
- Allows loop unrolling in HW
- Not limited to basic blocks (provided branch prediction)
- Lasting Contributions
 - Dynamic scheduling
 - Register renaming (in what way does the register name *change*?)
 - Load/store disambiguation

CSE 240A

Dean Tullsen

Scoreboard vs. Tomasulo, the score

	<u>Scoreboard</u>	<u>Tomasulo</u>
issue	when FU free	when RS free
read operands	from reg file	from reg file, CDB
write operands	to reg file	to CDB
structural hazards	functional units	reservation stations
WAW, WAR hazards	problem	no problem
register renaming	no	yes
instructions completing	no limit	1 / cycle (per CDB)
instructions beginning ex.	1 (per set of read ports)	no limit

CSE 240A

Dean Tullsen

Modern Architectures

- Alpha 21264+, MIPS R10K+, Pentium 4 use an *instruction queue*.
- They use explicit register renaming. Registers are not read until instruction issues (begins execution). Register renaming ensures no conflicts.

Div R5, R4, R2
Add R7, R5, R1
Sub R5, R3, R2
Lw R7, 1000(R5)



R1	PR23
R2	PR2
R3	PR17
R4	PR45
R5	PR13
R6	PR20
R7	PR30
...	

Dean Tullsen

CSE 240A

Modern Architectures

- Alpha 21264+, MIPS R10K+, Pentium 4 use an *instruction queue*.
- Uses explicit register renaming. Registers are not read until instruction issues (begins execution). Register renaming ensures no conflicts.

Div R5, R4, R2
Add R7, R5, R1
Sub R5, R3, R2
Lw R7, 1000(R5)



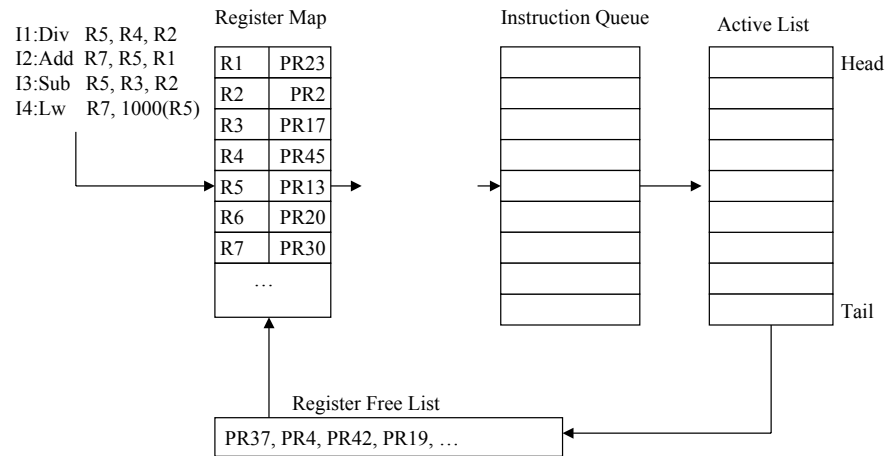
Div PR37, PR45, PR2
Add PR4, PR37, PR23
Sub PR42, PR17, PR2
Lw PR19, 1000(PR42)

R1	PR23
R2	PR2
R3	PR17
R4	PR45
R5	PR42
R6	PR20
R7	PR19
...	

Dean Tullsen

CSE 240A

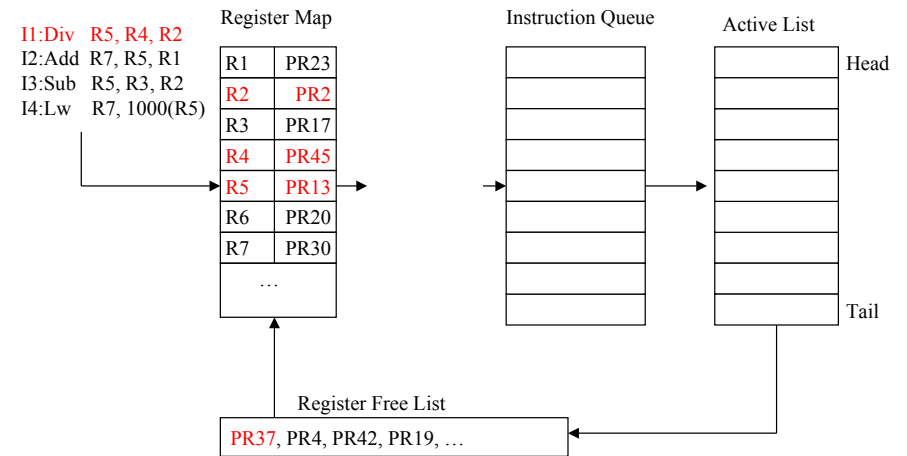
MIPS R10000, some detail



CSE 240A

Dean Tullsen

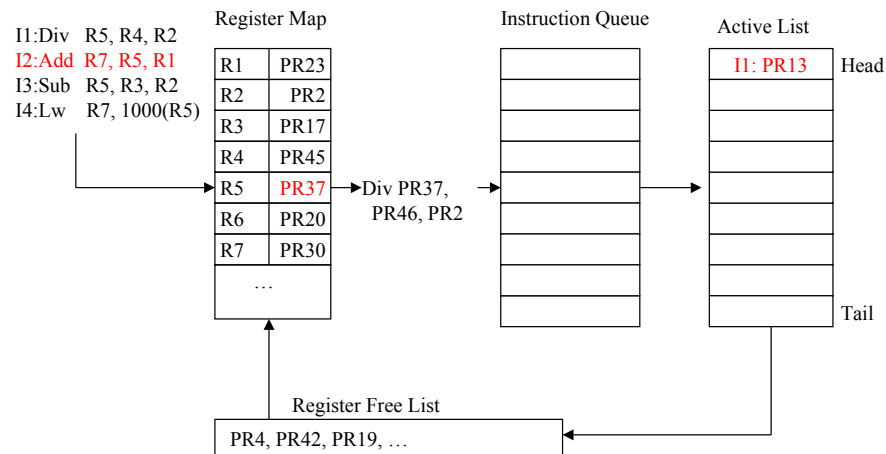
MIPS R10000, some detail



CSE 240A

Dean Tullsen

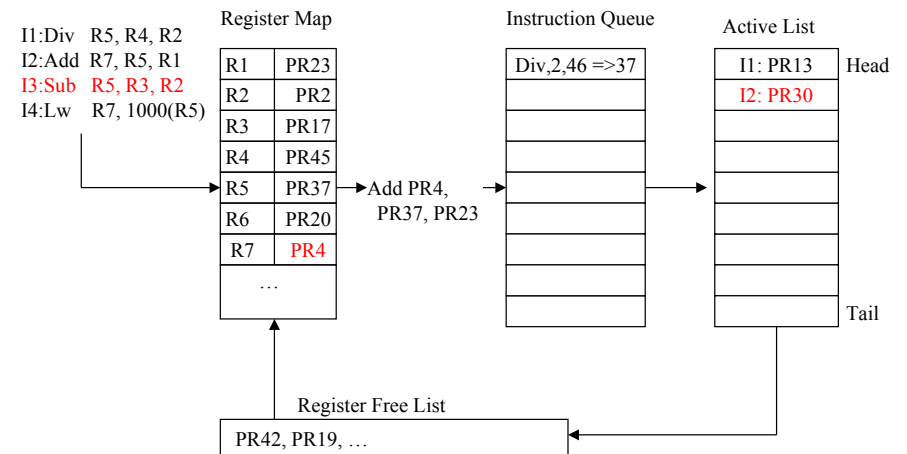
MIPS R10000, some detail



CSE 240A

Dean Tullsen

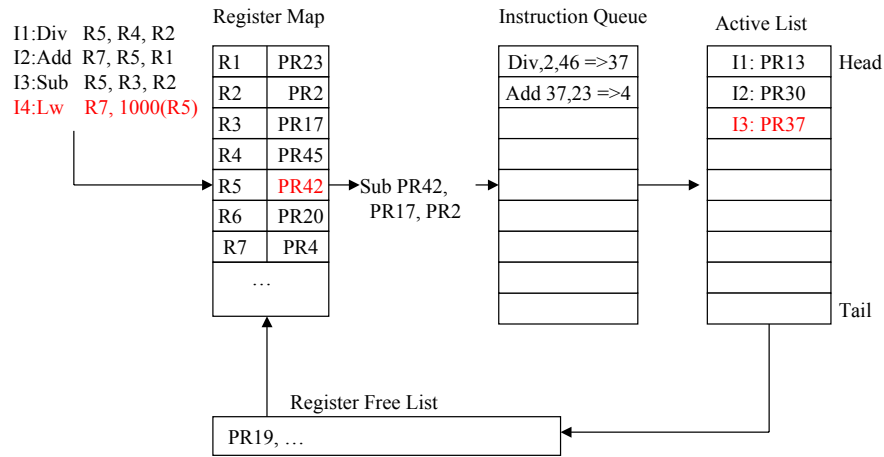
MIPS R10000, some detail



CSE 240A

Dean Tullsen

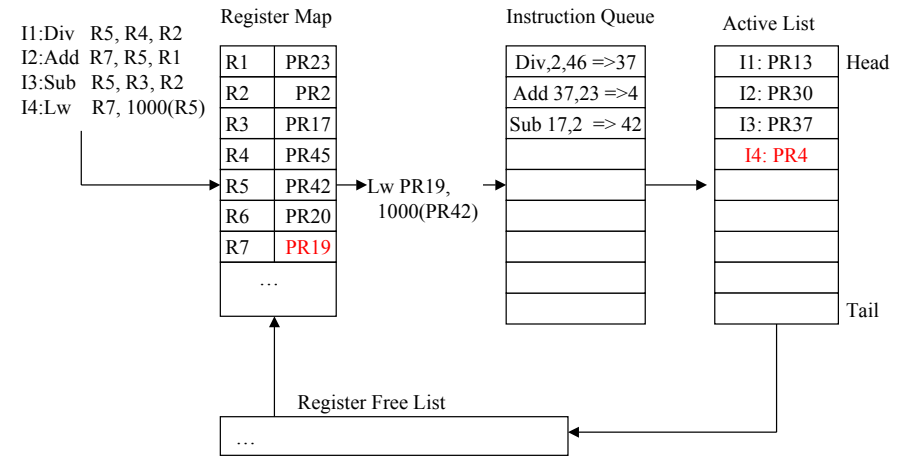
MIPS R10000, some detail



CSE 240A

Dean Tullsen

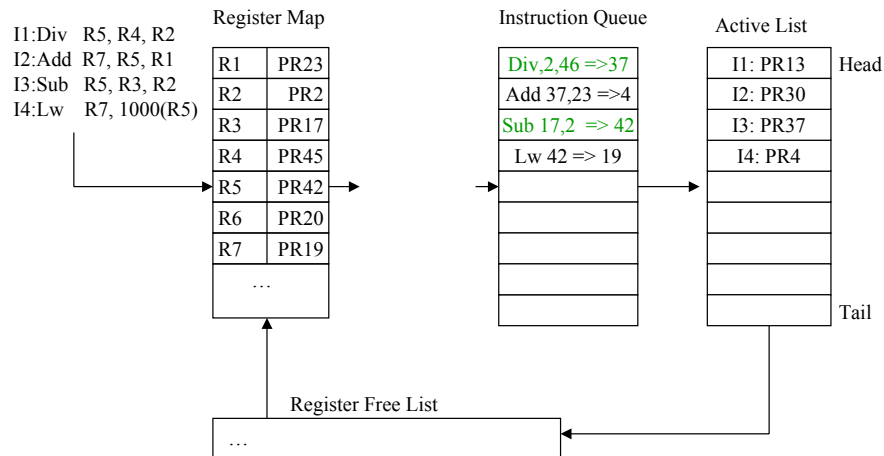
MIPS R10000, some detail



CSE 240A

Dean Tullsen

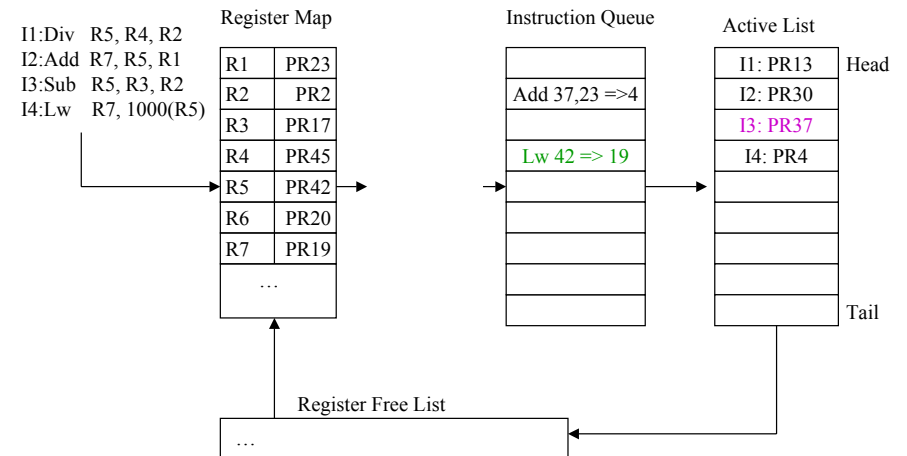
MIPS R10000, some detail



CSE 240A

Dean Tullsen

MIPS R10000, some detail

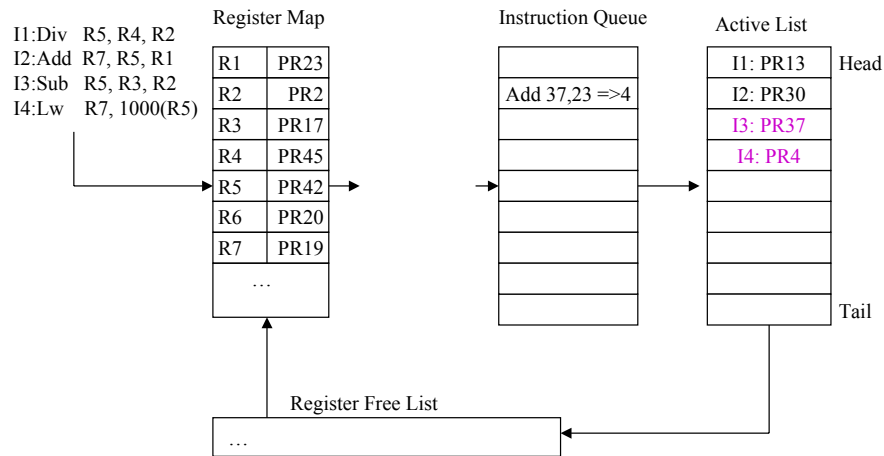


CSE 240A

Dean Tullsen

I3, producing register 42, completes, broadcasts a completion signal to IQ

MIPS R10000, some detail

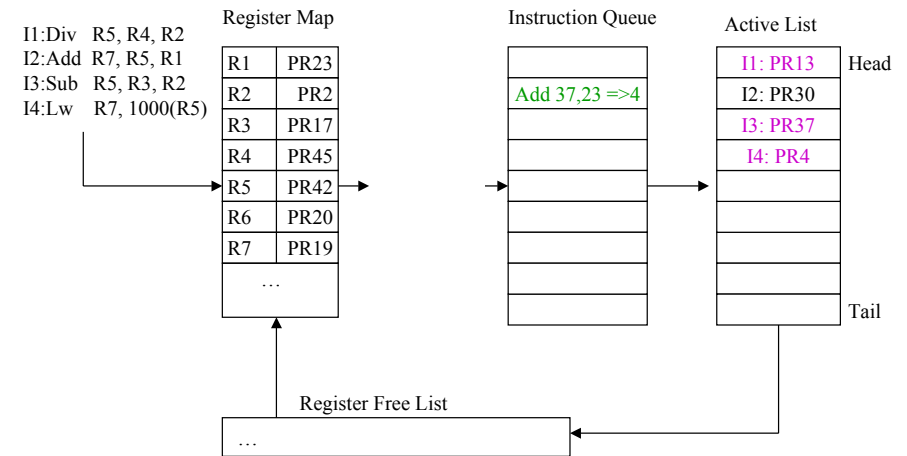


CSE 240A

I4, producing register 19, completes, broadcasts a completion signal to IQ

Dean Tullsen

MIPS R10000, some detail

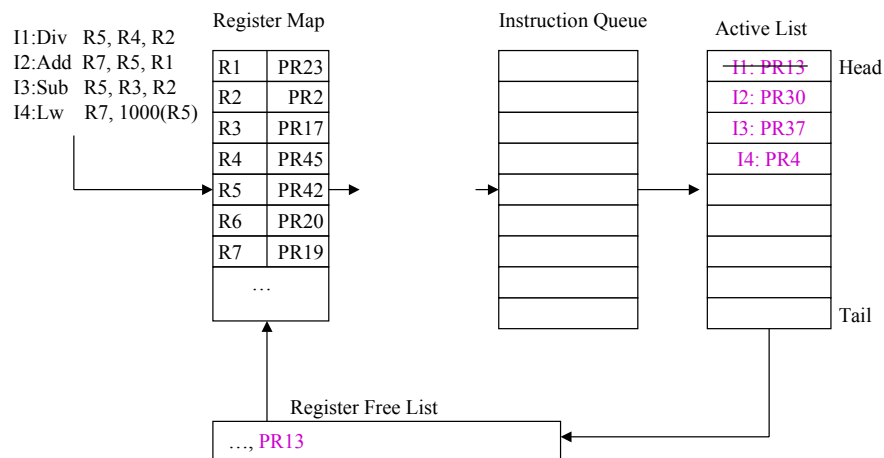


CSE 240A

I1, producing register 37, completes, broadcasts a completion signal to IQ

Dean Tullsen

MIPS R10000, some detail



CSE 240A

I2, producing register 4, completes, broadcasts a completion signal to IQ
I1 commits.

Dean Tullsen

Dynamic Scheduling Key Points

- Dynamic scheduling is code motion in HW.
- Dynamic scheduling can do things SW scheduling (static scheduling) cannot.
- Scoreboard, Tomasulo have various tradeoffs
- Register renaming eliminates WAW, WAR dependencies.
- To get cross-iteration parallelism, we need to eliminate WAW, WAR dependencies.

CSE 240A

Dean Tullsen